

Sum Yi Li, 505146702, sammyli0106@gmail.com  
Insiya Bambot, 605006720, insiyab8@gmail.com

## 1. Identity Block Output

```
Result = [0.8278597  0.10646105 2.1855583  0.56990707 0.          1.6353136
 1.454914    1.1312441 ]
```


## 2. Convolutional Block Output

```
Result = [0.          1.3407363 1.4184607 0.          0.          2.2878466
 0.8547394  0.5266884]
```

## 3. Accuracy of ResNet on CIFAR10 Test Set

```
Test loss: 0.9911629731178284
Test accuracy: 0.6735
```

## 4. Accuracy of training and validation plots



## 5. Is your model over-fitting or under-fitting?

From training our model:

```
Epoch 50/50  
25/25 [=====] - 45s 2s/step - loss: 0.6248 - acc:  
0.7780 - val_loss: 0.9912 - val_acc: 0.6735
```

Our training accuracy ends up being around 78% while our test accuracy was 67%. This indicates that our model overfits to the training data because the training accuracy is higher than the test accuracy by a significant amount (approximately a 10% difference). To address this, we could introduce dropout layers to our ResNet so that the trained model learns to use all the input data and not rely on a few repeated features in the training set.